

Few-shot Plant Disease Classification

Siddharth Khandelwal
190070062

Outline

- Motivation
- Literature Survey
- Datasets
- Few-shot learning
- CLIP
- Prompt Learning
- Implementation
- Experiments
- Conclusion
- References

Motivation

In a real world scenario one might not have many examples of classes of plant diseases on fields due to various factors :

1. Less documentation
2. Rare diseases

Few shot learning might prove helpful for such cases, to model and classify such few sampled diseases

Literature Survey

Bacterial Wilt Disease Detection

- Attacks more than 200 species of crops
- BWENet for grading severity of bacterial wilt disease
- Multi-spectral and Multi fractal analysis

Multi-modal approaches

- Image-text modality
- Hyperspectral imaging
- Soil moisture, temperature, rainfall

Datasets

Plant Village Dataset

- 256 x 256 size images
- 38 Plant leaf classes
- 54000 total images



Corn Northern Leaf Blight



Grape Leaf blight



Peach Bacterial_spot

Source : (<https://data.mendeley.com/datasets/tywbtsjrjv/1>)

Plant Doc Dataset

- Varying image sizes
- 28 Plant leaf classes
- 2600 total images



Bell pepper Leaf Spot



Tomato Early Blight



Healthy Apple leaf

Source : (<https://github.com/pratikkayal/PlantDoc-Dataset>)

Few-shot learning

Classify new data when you have only a few training samples with supervised information

N-way-K-shot classification :

- **Training set :**
Consists of N class labels each with K labelled images (where K can be around 5 to 10)
- **Query set :**
Contains the query images which need to be classified among the N classes

Possible approach : Gain experience from other similar problems. To this end, most approaches characterize few-shot learning as a **meta-learning** problem.

Meta Learning

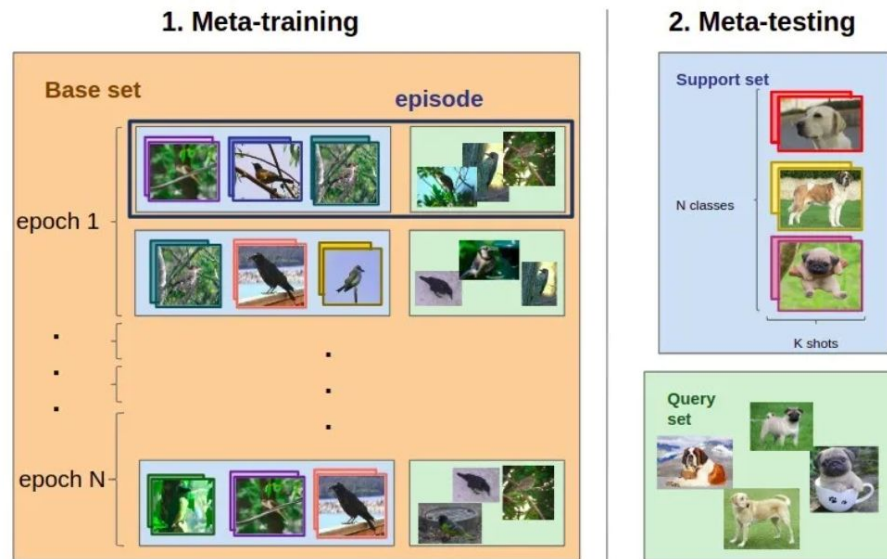
We treat our N-way-K-shot classification problem as the TEST data and a similar large base dataset as the meta learning training set (TRAIN)

Training:

N classes and K support images per each classes along with Q query images from the large TRAIN set

Test:

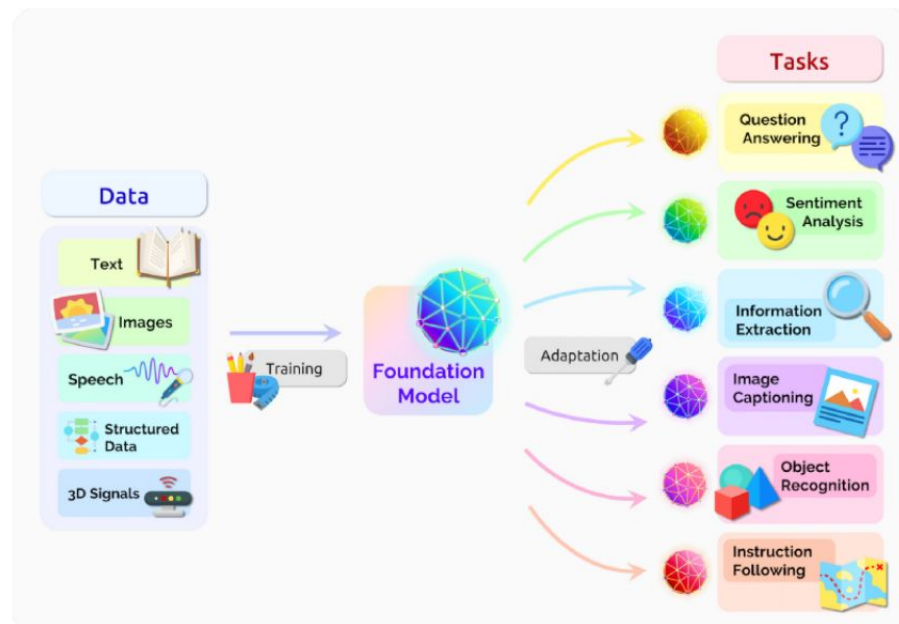
Classify our less labelled dataset using this trained model in similar episodes.



Meta learning [1]

Foundation Models in AI

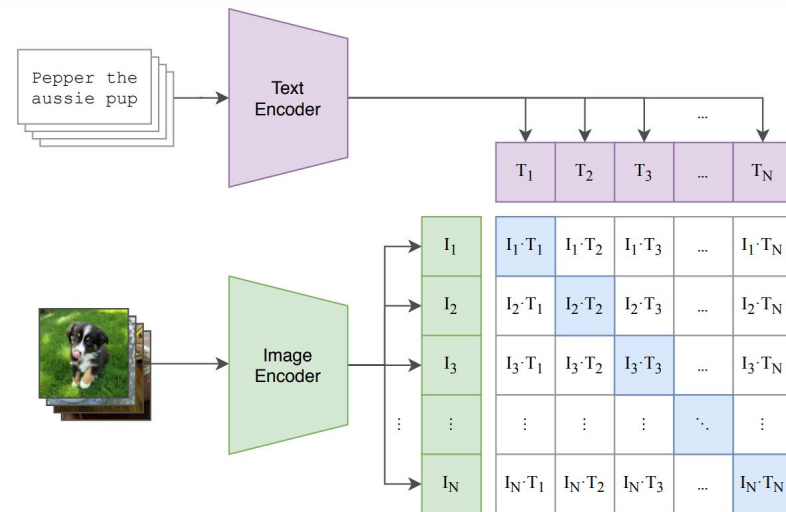
- Models that are trained on a broad set of unlabeled data
- Can be adapted for different tasks, with minimal fine-tuning
- Various application using self-supervised learning and transfer learning
- Eg: BERT, GPT-3, DALL-E 2, CLIP etc



Foundation Models in AI [2]

CLIP (Contrastive Language-Image Pre-training)

- Trained on wide variety of images and corresponding text from the internet
- During pre-training the model takes in a batch of N pairs of (image,text) input and learns a multi-modal embedding space by jointly training an image encoder and a text encoder
- Maximize the cosine similarity of the true N pairs and minimizing those of the rest possible pairs



Clip Model Pretraining [3]

Zero-shot with CLIP

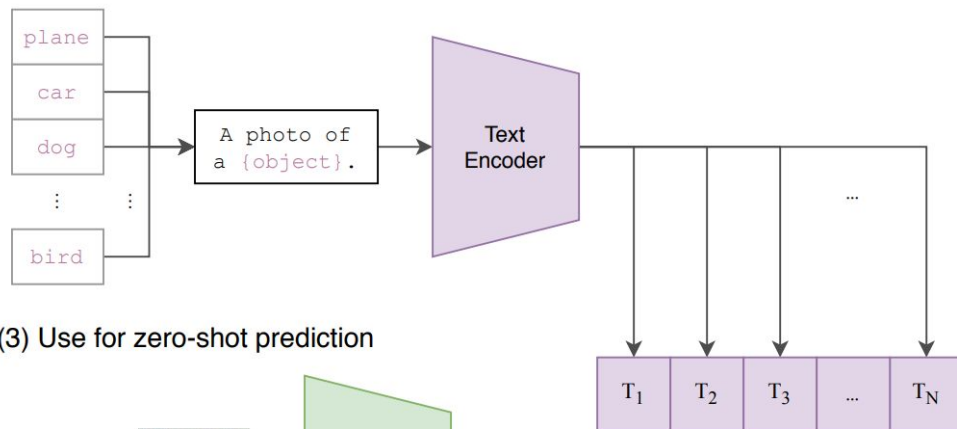
Aim : To classify an unseen image into infected or healthy using CLIP without training on Plant Village

Method : Visual classification by providing the names of the visual categories to be recognized to the pre-trained CLIP


Divided Plant Village Dataset into healthy and infected leaf classes

Result : **85 %** accuracy

(2) Create dataset classifier from label text



(3) Use for zero-shot prediction



Zero-shot with CLIP [3]

Problems with Zero shot using CLIP

- Prompt engineering is difficult
- Small changes in prompt result in large accuracy changes
- Cannot classify diseases

Need an automatic way to decide prompts for classes for better classification accuracy

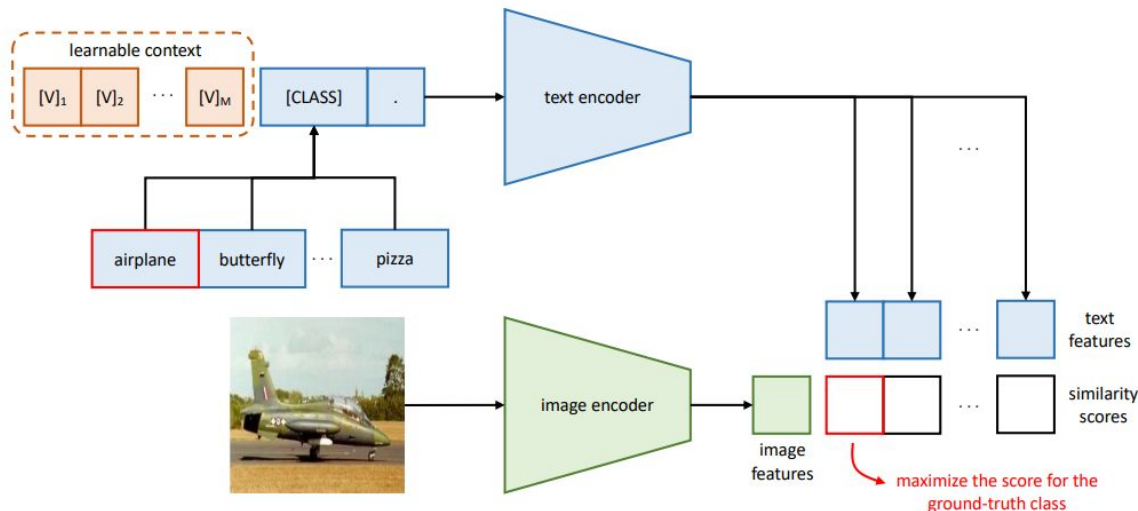
- CoOp
- CoCoOp

Prompt Learning

A) Context Optimization (CoOp Model)

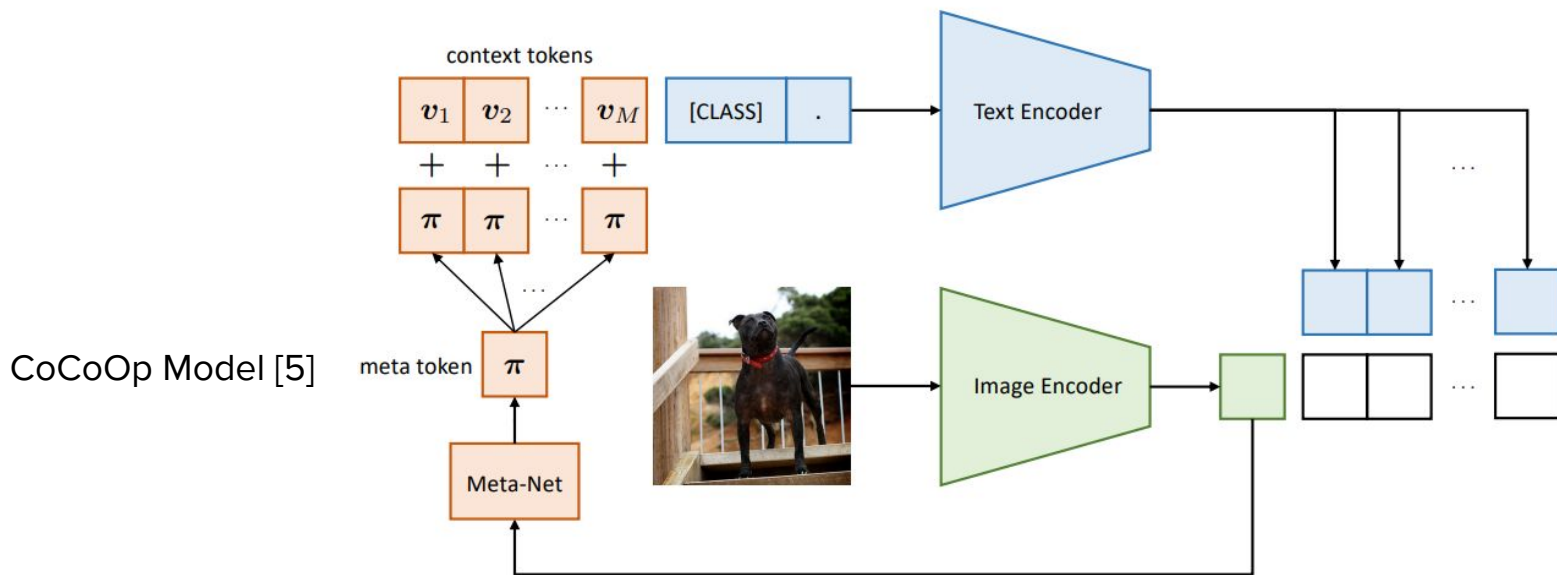
- Models a prompt's context words with learnable vectors
- Learning context through back propagation of classification loss
- Context learnt is static

CoOp Model [4]



B) Conditional Context Optimization (CoCoOp)

- Dynamic prompts
- Neural network on each instance (image) which adds information to context
- Better generalizability on unseen classes than CoOp



Implementation


Python framework : Pytorch

System : NVIDIA GeForce GTX 1080 Ti of VIP Lab

Experiments

Prototypical Few Shot Classification on Plant Doc Dataset

We see a drastic decrease in accuracy due to more complex images / noisy backgrounds etc in Plant Doc



(a) Few-shot

	Plant Village Dataset		Plant Doc Dataset	
	Training Accuracy (%)	Test Accuracy (%)	Training Accuracy (%)	Test Accuracy (%)
5-way-5-shot	93	82.8	82	46
5-way-10-shot	96	87	85	48

Experiments Contd.

Linear Probing on CLIP

- Logistic regression of features obtained from CLIP's image encoder
- Dataset : Plant Village
- Training & Validation : 38 classes having N labelled examples each (N shot)
- Test : 38 classes with around 300 examples each

	1 shot	2 shot	4 shot	8 shot	16 shot
Accuracy (%)	58.8 %	76.3 %	87 %	91.4 %	94.2 %

Note: The test set and training set have common classes between them in this case

Experiments Contd.

Few shot learning using CoOp on Plant Village

- Dataset : Plant Village
- Training & Validation : 38 classes having K labelled examples each
- Test : 38 classes with around 300 examples each

	1 shot	2 shot	4 shot	8 shot	16 shot
Accuracy (%)	27 \pm 3	50	53 \pm 5	74 \pm 2	82

Note: The test set and training set have common classes between them in this case

Experiments Contd.

Generalization on Unseen New Classes

- CoOp and CoCoOp are compared
- Dataset : Plant Village
- Training & Validation set : 30 classes with around 1000 examples each
- Test set : 8 new classes with 20 examples each

	CoOp		CoCoOp	
	Training Accuracy (%)	Test Accuracy (%)	Training Accuracy (%)	Test Accuracy (%)
5-shot	78	26	58	38
10-shot	86	13	66	20

Conclusion

- Learning with limited examples per class with Linear Probe and using CoOp feasible
- CoCoOp and CoOp models does not give good results on Plant Village for generalization on unseen classes

References

- [1] E. Bennequin, “Few-Shot Image Classification with Meta-learning”,2018
- [2] Bommasani, R., Hudson, D., Adeli, E., Altman, R., & Arora et. al. (2021). “On the Opportunities and Risks of Foundation Models”
- [3] Radford, A., Kim, J., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., Sastry, G., Askell, A., Mishkin, P., Clark, J., Krueger, G., & Sutskever, I.. (2021). Learning Transferable Visual Models From Natural Language Supervision.
- [4] Kaiyang Zhou, Jingkang Yang, Chen Change Loy, & Ziwei Liu (2022). Learning to Prompt for Vision-Language Models. *International Journal of Computer Vision*, 130(9), 2337–2348.
- [5] Zhou, K., Yang, J., Loy, C., & Liu, Z.. (2022). Conditional Prompt Learning for Vision-Language Models

Thank You